

Model Development Phase Template

Date	11-07-2024
Team ID	SWTID1720433291
Project Title	CovidVision: Advanced COVID-19 Detection From Lung X-Rays With Deep Learning
Maximum Marks	10 Marks

Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be showcased in the future through a screenshot. The model validation and evaluation report will include a summary and training and validation performance metrics for multiple models, presented through respective screenshots.

Initial Model Training Code (5 marks):

CNN:

```
[4] train_data_dir= "Xray_train_data"

[5] IMAGE_SIZE = (256, 256)
    IMAGE_SHAPE = IMAGE_SIZE + (3,)

data_augmentation = keras.Sequential([
    tf.keras.layers.RandomRotation(
        factor=(-0.2, 0.3),
        fill_mode='reflect',
        interpolation='bilinear',
        seed=None
    ),
    tf.keras.layers.Rescaling(
        scale=1/.255,
        offset=0.0
    ),
])

[6]
```

```
model=tf.keras.Sequential([
    data_augmentation,
    tf.keras.layers.Conv2D(8, (3,3), activation='relu', input_shape=IMAGE_SHAPE),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2), strides=None, padding="valid"),
    tf.keras.layers.Conv2D(16, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2), strides=None, padding="valid"),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(32),
    tf.keras.layers.Dropout(.2, input_shape=(32,)),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

[7]

```
from tensorflow.keras.optimizers import Adam

model.compile(optimizer=Adam(learning_rate=0.0001),
              loss='binary_crossentropy',
              metrics = ['accuracy'])
```

[42]

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
img_height, img_width=64,64
batch_size=16
train_datagen = ImageDataGenerator(validation_split=0.3) # set validation split

train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='binary',
    subset='training') # set as training data

# Splitting images for validation set
validation_generator = train_datagen.flow_from_directory(
    train_data_dir, # same directory as training data
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='binary',
    subset='validation') # set as validation data
```

[43]

... Found 14729 images belonging to 2 classes.
Found 6311 images belonging to 2 classes.

XCEPTION:

```
train_datagen = ImageDataGenerator(preprocessing_function=tf.keras.applications.xception.preprocess_input, zoom_range=0.1, brightness_range=[0.5, 1.3],
                                  width_shift_range=0.1, height_shift_range=0.1, validation_split=0.1)
test_datagen=ImageDataGenerator(preprocessing_function=tf.keras.applications.xception.preprocess_input)
```

```
BATCH_SIZE=64
path="Coronahack-Chest-XRay-Dataset/Coronahack-Chest-XRay-Dataset/"
train_images=train_datagen.flow_from_dataframe(dataframe=df_meta[df_meta["Dataset_type"]=="TRAIN"],x_col='X_ray_image_name',y_col='Label',
                                                color_mode='rgb', class_mode='categorical',target_size=(150,150),
                                                batch_size=BATCH_SIZE,shuffle=True,seed=1234,subset='training',
                                                directory=path+"/train")

val_images=train_datagen.flow_from_dataframe(dataframe=df_meta[df_meta["Dataset_type"]=="TRAIN"],x_col='X_ray_image_name',y_col='Label',color_mode='rgb',
                                                class_mode='categorical',target_size=(150,150),batch_size=BATCH_SIZE,shuffle=True,seed=1234,
                                                subset='validation', directory=path+"/train")

test_images = test_datagen.flow_from_dataframe(dataframe=df_meta[df_meta["Dataset_type"]=="TEST"],x_col='X_ray_image_name',y_col='Label',
                                                color_mode='rgb',class_mode='categorical',target_size=(150,150),batch_size=BATCH_SIZE,shuffle=False,
                                                directory=path+"/test")
```

```
inputs = tf.keras.layers.Input((150,150,3))
base_model=tf.keras.applications.xception.Xception(include_top=False, weights="imagenet",input_shape=(150,150,3), pooling='avg')
x=base_model(inputs)
output=layers.Dense(2, activation='sigmoid')(x)
model=tf.keras.models.Model(inputs=inputs, outputs=output)

model.compile(Adamax(learning_rate=1e-4), loss='binary_crossentropy',metrics=['accuracy'])
history = model.fit(train_images, validation_data=val_images, epochs=30)
```

Model Validation and Evaluation Report (5 marks):

Model	Summary	Training and Validation Performance Metrics																														
CNN	 <p>Model: "sequential_1"</p> <table border="1"> <thead> <tr> <th>Layer (type)</th> <th>Output Shape</th> <th>Param #</th> </tr> </thead> <tbody> <tr> <td>Sequential (Conv2D_1)</td> <td>3</td> <td>1 (Conv2D1)</td> </tr> <tr> <td>Sequential (Conv2D_2)</td> <td>3</td> <td>1 (Conv2D2)</td> </tr> <tr> <td>Max pooling2D (MaxPooling2D)</td> <td>3</td> <td>0 (Conv2D1)</td> </tr> <tr> <td>conv2d_2 (Conv2D)</td> <td>3</td> <td>1 (Conv2D1)</td> </tr> <tr> <td>Max pooling2D_2 (MaxPooling2D)</td> <td>3</td> <td>0 (Conv2D1)</td> </tr> <tr> <td>Flatten (Flatten)</td> <td>3</td> <td>0 (Conv2D1)</td> </tr> <tr> <td>Dense (Dense)</td> <td>3</td> <td>1 (Conv2D1)</td> </tr> <tr> <td>Dense_1 (Dense)</td> <td>3</td> <td>1 (Conv2D1)</td> </tr> <tr> <td>Dense_2 (Dense)</td> <td>3</td> <td>1 (Conv2D1)</td> </tr> </tbody> </table> <p>Total params: 1 (0.00 s) Trainable params: 1 (0.00 s) Non-trainable params: 0 (0.00 s)</p>	Layer (type)	Output Shape	Param #	Sequential (Conv2D_1)	3	1 (Conv2D1)	Sequential (Conv2D_2)	3	1 (Conv2D2)	Max pooling2D (MaxPooling2D)	3	0 (Conv2D1)	conv2d_2 (Conv2D)	3	1 (Conv2D1)	Max pooling2D_2 (MaxPooling2D)	3	0 (Conv2D1)	Flatten (Flatten)	3	0 (Conv2D1)	Dense (Dense)	3	1 (Conv2D1)	Dense_1 (Dense)	3	1 (Conv2D1)	Dense_2 (Dense)	3	1 (Conv2D1)	 <pre> Trainer: model_fit(train_generator,validation_data=validation_generator,epochs=30) Epoch 1/25: 17s 14ms/step - accuracy: 0.7880 - loss: 0.4094 - val_accuracy: 0.4351 - val_loss: 1.1448 Epoch 2/25: 18s 10ms/step - accuracy: 0.7833 - loss: 0.4078 - val_accuracy: 0.4345 - val_loss: 1.1631 Epoch 3/25: 18s 11ms/step - accuracy: 0.7890 - loss: 0.4174 - val_accuracy: 0.4246 - val_loss: 1.1738 Epoch 4/25: 18s 11ms/step - accuracy: 0.7887 - loss: 0.4169 - val_accuracy: 0.4238 - val_loss: 1.1882 Epoch 5/25: 18s 10ms/step - accuracy: 0.7839 - loss: 0.4087 - val_accuracy: 0.4253 - val_loss: 1.1876 Epoch 6/25: 18s 10ms/step - accuracy: 0.7834 - loss: 0.4039 - val_accuracy: 0.4258 - val_loss: 1.1921 Epoch 7/25: 18s 11ms/step - accuracy: 0.7889 - loss: 0.4192 - val_accuracy: 0.4256 - val_loss: 1.2632 Epoch 8/25: 18s 10ms/step - accuracy: 0.7885 - loss: 0.4192 - val_accuracy: 0.4262 - val_loss: 1.2132 Epoch 9/25: 18s 10ms/step - accuracy: 0.7888 - loss: 0.4185 - val_accuracy: 0.4253 - val_loss: 1.2234 Epoch 10/25: 18s 11ms/step - accuracy: 0.7833 - loss: 0.4086 - val_accuracy: 0.4258 - val_loss: 1.2258 Epoch 11/25: 18s 10ms/step - accuracy: 0.7890 - loss: 0.4083 - val_accuracy: 0.4238 - val_loss: 1.2257 Epoch 12/25: 18s 10ms/step - accuracy: 0.7892 - loss: 0.4153 - val_accuracy: 0.4256 - val_loss: 1.2488 Epoch 13/25: 18s 10ms/step - accuracy: 0.7890 - loss: 0.4152 - val_accuracy: 0.4264 - val_loss: 1.2436 Epoch 14/25: 18s 10ms/step - accuracy: 0.7888 - loss: 0.4013 - val_accuracy: 0.4253 - val_loss: 1.2526 Epoch 15/25: 18s 11ms/step - accuracy: 0.7836 - loss: 0.4165 - val_accuracy: 0.4264 - val_loss: 1.2623 Epoch 16/25: 18s 11ms/step - accuracy: 0.7888 - loss: 0.4081 - val_accuracy: 0.4266 - val_loss: 1.2727 Epoch 17/25: 18s 10ms/step - accuracy: 0.7833 - loss: 0.4173 - val_accuracy: 0.4253 - val_loss: 1.2834 Epoch 18/25: 18s 10ms/step - accuracy: 0.7883 - loss: 0.4146 - val_accuracy: 0.4258 - val_loss: 1.2876 Epoch 19/25: 18s 11ms/step - accuracy: 0.7890 - loss: 0.4126 - val_accuracy: 0.4262 - val_loss: 1.2938 Epoch 20/25: 18s 11ms/step - accuracy: 0.7889 - loss: 0.4139 - val_accuracy: 0.4239 - val_loss: 1.3131 Epoch 21/25: 18s 11ms/step - accuracy: 0.7893 - loss: 0.4178 - val_accuracy: 0.4253 - val_loss: 1.3155 Epoch 22/25: 18s 11ms/step - accuracy: 0.7890 - loss: 0.4087 - val_accuracy: 0.4259 - val_loss: 1.3172 Epoch 23/25: 18s 11ms/step - accuracy: 0.7893 - loss: 0.4038 - val_accuracy: 0.4253 - val_loss: 1.3142 Epoch 24/25: 18s 11ms/step - accuracy: 0.7890 - loss: 0.4089 - val_accuracy: 0.4256 - val_loss: 1.3423 Epoch 25/25: 18s 10ms/step - accuracy: 0.7889 - loss: 0.4123 - val_accuracy: 0.4259 - val_loss: 1.3523 </pre>
Layer (type)	Output Shape	Param #																														
Sequential (Conv2D_1)	3	1 (Conv2D1)																														
Sequential (Conv2D_2)	3	1 (Conv2D2)																														
Max pooling2D (MaxPooling2D)	3	0 (Conv2D1)																														
conv2d_2 (Conv2D)	3	1 (Conv2D1)																														
Max pooling2D_2 (MaxPooling2D)	3	0 (Conv2D1)																														
Flatten (Flatten)	3	0 (Conv2D1)																														
Dense (Dense)	3	1 (Conv2D1)																														
Dense_1 (Dense)	3	1 (Conv2D1)																														
Dense_2 (Dense)	3	1 (Conv2D1)																														

XCEPTION

```

[1]: win1.com>D
Model: "Xception"

```

Layer (type)	Action/Shape	Param #
conv2d_1 (Conv2D)	convolve, (3, 3, 3, 32)	
maxpooling2d (MaxPooling2D)	convolve, (2, 2)	16,384 (8)
conv2d_2 (Conv2D)	convolve, (3, 3)	6,554

```

Total params: 1,407,206 (270.19 MB)
Trainable params: 17,43,342 (336.09 MB)
Non-trainable params: 16,384 (8)
Optimizer params: 17,43,342 (336.09 MB)

```

```

[5]: model.compile(optimizer=learning_rate_scheduler(lr_scheduler), loss='binary_crossentropy', metrics=['accuracy'])
library = model.compile(optimizer=learning_rate_scheduler(lr_scheduler), loss='binary_crossentropy', metrics=['accuracy'])

Epoch 1/50
Epoch 1/50: 312s 4s/step - accuracy: 0.8558 - loss: 0.4613 - val_accuracy: 0.4754 - val_loss: 1.6222
Epoch 2/50: 318s 4s/step - accuracy: 0.8563 - loss: 0.1374 - val_accuracy: 0.8333 - val_loss: 0.3585
Epoch 3/50: 318s 4s/step - accuracy: 0.8785 - loss: 0.8038 - val_accuracy: 0.9242 - val_loss: 0.1594
Epoch 4/50: 318s 4s/step - accuracy: 0.8888 - loss: 0.8637 - val_accuracy: 0.9478 - val_loss: 0.1115
Epoch 5/50: 318s 4s/step - accuracy: 0.8781 - loss: 0.8711 - val_accuracy: 0.9394 - val_loss: 0.1662
Epoch 6/50: 318s 4s/step - accuracy: 0.8884 - loss: 0.8811 - val_accuracy: 0.9478 - val_loss: 0.1117
Epoch 7/50: 318s 4s/step - accuracy: 0.8877 - loss: 0.8427 - val_accuracy: 0.9754 - val_loss: 0.0884
Epoch 8/50: 318s 4s/step - accuracy: 0.8889 - loss: 0.8363 - val_accuracy: 0.9775 - val_loss: 0.0708
Epoch 9/50: 318s 4s/step - accuracy: 0.8883 - loss: 0.8332 - val_accuracy: 0.9754 - val_loss: 0.1185
Epoch 10/50: 318s 4s/step - accuracy: 0.8885 - loss: 0.8326 - val_accuracy: 0.9811 - val_loss: 0.0855
Epoch 11/50: 318s 4s/step - accuracy: 0.8882 - loss: 0.8366 - val_accuracy: 0.9432 - val_loss: 0.1684
Epoch 12/50: 318s 4s/step - accuracy: 0.8854 - loss: 0.8219 - val_accuracy: 0.9662 - val_loss: 0.1845
Epoch 13/50: 318s 4s/step - accuracy: 0.8864 - loss: 0.8488 - val_accuracy: 0.9678 - val_loss: 0.0796
Epoch 14/50: 318s 4s/step - accuracy: 0.8923 - loss: 0.8238 - val_accuracy: 0.9659 - val_loss: 0.0749
Epoch 15/50: 318s 4s/step - accuracy: 0.8941 - loss: 0.8227 - val_accuracy: 0.9722 - val_loss: 0.1267
Epoch 16/50: 318s 4s/step - accuracy: 0.8925 - loss: 0.8197 - val_accuracy: 0.9678 - val_loss: 0.1785
Epoch 17/50: 318s 4s/step - accuracy: 0.8946 - loss: 0.8118 - val_accuracy: 0.9384 - val_loss: 0.1112
Epoch 18/50: 318s 4s/step - accuracy: 0.8989 - loss: 0.8278 - val_accuracy: 0.9848 - val_loss: 0.0428
Epoch 19/50: 318s 4s/step - accuracy: 0.8994 - loss: 0.8161 - val_accuracy: 0.9811 - val_loss: 0.0638
Epoch 20/50: 318s 4s/step - accuracy: 0.8933 - loss: 0.8186 - val_accuracy: 0.9716 - val_loss: 0.0817
Epoch 21/50: 318s 4s/step - accuracy: 0.8956 - loss: 0.8126 - val_accuracy: 0.9754 - val_loss: 0.0706
Epoch 22/50: 318s 4s/step - accuracy: 0.8952 - loss: 0.8163 - val_accuracy: 0.9807 - val_loss: 0.0703
Epoch 23/50: 318s 4s/step - accuracy: 0.8953 - loss: 0.8118 - val_accuracy: 0.9545 - val_loss: 0.1112
Epoch 24/50: 318s 4s/step - accuracy: 0.8971 - loss: 0.8087 - val_accuracy: 0.9648 - val_loss: 0.0864
Epoch 25/50: 318s 4s/step - accuracy: 0.8971 - loss: 0.8082 - val_accuracy: 0.9621 - val_loss: 0.1288
Epoch 26/50: 318s 4s/step - accuracy: 0.8985 - loss: 0.8066 - val_accuracy: 0.9662 - val_loss: 0.1842
Epoch 27/50: 318s 4s/step - accuracy: 0.8936 - loss: 0.8166 - val_accuracy: 0.9716 - val_loss: 0.0728
Epoch 28/50: 318s 4s/step - accuracy: 0.8958 - loss: 0.8125 - val_accuracy: 0.9792 - val_loss: 0.0564
Epoch 29/50: 318s 4s/step - accuracy: 0.8982 - loss: 0.8118 - val_accuracy: 0.9775 - val_loss: 0.0851
Epoch 30/50: 318s 4s/step - accuracy: 0.8978 - loss: 0.8082 - val_accuracy: 0.9545 - val_loss: 0.1804

model.save("xception_model.keras")

```